

O2 FEB 2023 – HPC House Power Consumption

HADOOP & MAPREDUCE

```
public class CityCount implements org.apache.hadoop.io.Writable {
    public String city;
    public int count;

    public String getCity() {
        return city;
    }
    public void setCity(String city) {
        this.city = city;
    }
    public int getCount() {
        return count;
    }
    public void setCount(int count) {
        this.count = count;
    }

    public String toString() {
        return city + "," + count;
    }
    @Override
    public void write(DataOutput out) throws IOException {
        out.writeUTF(city);
        out.writeInt(count);
    }

    @Override
    public void readFields(DataInput in) throws IOException {
        city = in.readUTF();
        count = in.readInt();
    }
}

class Mapper1BigData extends Mapper<
    LongWritable,
    Text,
    Text,
    IntWritable> {

    protected void map(
```

```

        LongWritable key,
        Text value,
        Context context) throws IOException, InterruptedException {

    String[] fields = value.toString().split(",");
    String city = fields[1];
    Double sqm = Double.parseDouble(fields[3]);

    if(sqm < 60) {
        context.write(new Text(city), new IntWritable(1));
    }
}

}

class Reducer1BigData extends Reducer<Text, IntWritable, Text, IntWritable> {

    private String maxCity;
    private int maxCount;

    @Override
    protected void setup(Context context) throws IOException,
InterruptedException {
        this.maxCity = null;
        this.maxCount = -1;
    }

    @Override
    protected void reduce(
        Text key,
        Iterable<IntWritable> values,
        Context context) throws IOException, InterruptedException {
        int val = -1;
        int sum = 0;
        String city = key.toString();
        for (IntWritable v : values) {
            val = v.get();
            sum += val;
        }

        if (maxCount == -1 || sum > maxCount || (sum == maxCount &&
city.compareTo(maxCity) < 0)) {
            maxCount = sum;
            maxCity = city;
        }
    }
}

```

```

    }

    @Override
    protected void cleanup(Context context) throws IOException,
        InterruptedException {
        if (this.maxCity != null) {
            context.write(new Text(maxCity), new IntWritable(maxCount));
        }
    }
}

```

```

class Mapper2BigData extends Mapper<
    Text,
    Text,
    NullWritable,
    CityCount> {

    protected void map(
        Text key,
        Text value,
        Context context) throws IOException, InterruptedException {

        CityCount localCityCount = new CityCount();
        localCityCount.city=key.toString();
        localCityCount.count=Integer.parseInt(value.toString());

        context.write(NullWritable.get(), localCityCount);
    }
}

```

```

class Reducer2BigData extends Reducer<
    NullWritable,
    CityCount,
    Text,
    IntWritable> {

    @Override
    protected void reduce(
        NullWritable key,
        Iterable<CityCount> values,
        Context context) throws IOException, InterruptedException {

```

```
String maxCity=null;
int maxCount=-1;

for(CityCount v : values) {
    if(maxCount == -1 || v.count > maxCount ||
        (v.count == maxCount && v.city.compareTo(maxCity) < 0)) {
        maxCount = v.count;
        maxCity = v.city;
    }
}

context.write(new Text(maxCity), new IntWritable(maxCount));
}

}
```

SPARK

```
import pyspark

from pyspark import SparkContext
from pyspark.sql import SparkSession

spark = SparkSession.builder.getOrCreate()
sc = spark.sparkContext

housePath = "data/Houses.txt"
consumptionPath = "data/DailyPowerConsumption.txt"

outputPath1 = "outPart1/"
outputPath2 = "outPart2/"

houseRDD = sc.textFile(housePath)
consumptionRDD = sc.textFile(consumptionPath)

# Part 1
# filter only the readings associated with year 2022
consumption2022 = consumptionRDD.filter(lambda s:
s.split(",")[1].startswith("2022"))

# compute the total amount of energy consumed in year 2022 for each house
# key = houseID
# value = kWh consumed in year 2022

def mapCons(s):
    fields = s.split(",")
    hid = fields[0]
    consumption = float(fields[2])
    return (hid, consumption)

totalCons2022 = consumption2022.map(mapCons)\
                                .reduceByKey(lambda v1, v2: v1 + v2)

# compute the avg power consumption per day
# key = houseID
# value = avg kWh consumed per day in year 2022
# and filter only those with high avg consumption
highAvgDailyCons = totalCons2022\
                    .mapValues(lambda v: v / 365)\
                    .filter(lambda i: i[1] > 30)
```

```

# compute the pairRDD house -> country
# key = houseID
# value = country

def mapHouseCountry(s):
    fields = s.split(",")
    hid = fields[0]
    country = fields[2]

    return (hid, country)

houseCountry = houseRDD.map(mapHouseCountry)

# keep an RDD containing countries with at least one house with high
# avg power consumption
countriesWithHighAvgPwrConsHouses = houseCountry.join(highAvgDailyCons)\
    .map(lambda it: it[1][0]) # Country

# compute an RDD with all the countries
# and subtract the countries with at least one house with high avg power
# consumption
res1 = houseCountry.map(lambda v: v[1])\
    .distinct()\
    .subtract(countriesWithHighAvgPwrConsHouses)

res1.saveAsTextFile(outputPath1)

# Part 2
# keep only the houses for which the total power consumption over 2021 is > 10000
# kWh

def mapHidCons(s):
    fields = s.split(",")
    hid = fields[0]
    consumption = float(fields[2])

    return (hid, consumption)

highTotalPowerCons2021 = consumptionRDD\
    .filter(lambda s: s.split(",")[1].startswith("2021"))\
    .map(mapHidCons)\
    .reduceByKey(lambda v1, v2: v1 + v2)\
    .filter(lambda v: v[1]>10000)

```

```

# compute an RDD with
# key = houseID
# value = (country, city)

def mapHidCountryCity(s):
    fields = s.split(",")
    hid = fields[0]
    city = fields[1]
    country = fields[2]

    return (hid, (country, city))

citiesRDD = houseRDD.map(mapHidCountryCity)

# join the two RDDs and count for each city the number of houses with high annual
# power consumption
# and filter only those cities with value > 500
# key = (country, city)
# value = #houses with high power consumption
highPwrConsHousesPerCity = highTotalPowerCons2021.join(citiesRDD)\
    .map(lambda p: (p[1][1], 1))\
    .reduceByKey(lambda v1, v2: v1 + v2)\
    .filter(lambda p: p[1]>500)

# count for each country the number of cities with at least 500 houses with high
# annual power consumption

# Map each input pair to a new pair
# key = country
# value = +1 (one more city for this country with at least 500 houses with high
# annual power consumption)
countryOneMorehighPwrConsCity = highPwrConsHousesPerCity\
    .map(lambda p: (p[0][0], 1))

# Map each country to pair
# key = country
# value = 0
# This is used to keep also countries without cities with at least
# 500 houses with high annual power consumption
countriesZero = houseCountry.map(lambda p: (p[1], 0))

# count for each country the number of cities with at least 500 houses with high
# annual power consumption

```

```
# Union countryOneMorehighPwrConsCity with countriesZero and apply reduceByKey
# Output
# key = country
# value = number of cities with at least 500 houses with high power consumption
highPwrConsCitiesPerCountry = countryOneMorehighPwrConsCity.union(countriesZero)\
    .reduceByKey(lambda v1, v2: v1 + v2)

highPwrConsCitiesPerCountry.saveAsTextFile(outputPath2)
```