

06 JUNE 2022 – MONITOR DATA CENTERS

MAPREDUCE & HADOOP

```
class Mapper extends org.apache.hadoop.mapreduce.Mapper<
    LongWritable, // Input key type
    Text,          // Input value type
    Text,          // Output key type
    Text> { // Output value type

    protected void map(
        LongWritable key, // Input key type
        Text value,       // Input value type
        Context context) throws IOException, InterruptedException {

        String[] fields = value.toString().split(",");
        String year = fields[1].split("/")[0];
        String codDC = fields[0];
        double pwrCons = Double.parseDouble(fields[2]);

        if(year.equals("2020") || year.equals("2021"))
            if(pwrCons >= 1000)
                context.write(new Text(codDC), new Text(year));
    }
}
```

```
class Reducer extends org.apache.hadoop.mapreduce.Reducer<
    Text,          // Input key type
    Text,          // Input value type
    Text,          // Output key type
    NullWritable> { // Output value type

    @Override
    protected void reduce(
        Text key, // Input key type
        Iterable<Text> values, // Input value type
        Context context) throws IOException, InterruptedException {

        int count2020 = 0;
        int count2021 = 0;

        for(Text year : values) {
            String val = year.toString();
            if(val.equals("2020"))
                count2020++;
        }
    }
}
```

```
        else
            count2021++;
    }
    if(count2021 > count2020)
        context.write(key, NullWritable.get());
    }
}
```

SPARK

```
import pyspark

from pyspark import SparkContext
from pyspark.sql import SparkSession

spark = SparkSession.builder.getOrCreate()
sc = spark.sparkContext

# companiesPath = "data/Companies.txt" # Useless for this program
dataCenterPath = "data/DataCenters.txt"
dailyPwrConsPath = "data/DailyPowerConsumption.txt"

outputPath1 = "outPart1/"
outputPath2 = "outPart2/"

# Define the rdds associated with the used input files
# CodDC,CodC,City,Country,Continent
dataCenterRDD = sc.textFile(dataCenterPath).cache()
# CodDC,Date,kWh
pwrConsRDD = sc.textFile(dailyPwrConsPath)

#####
# PART 1
#####

# Count the total number of data centers worl-wide and compute the threshold (90%
of the data centers)
threshold = int(dataCenterRDD.count() * 0.9)

# Filter only the dates associated with high power consumption (>=1000 KWh)
#
# Prepare a pairRDD with
# key = date
# value = +1
#
# Finally, to count the number of data centers with high power consumption for
each day/date.
highPwrConsDCPerDay = pwrConsRDD.filter(lambda line: float(line.split(",")[2]) >=
1000)\
    .map(lambda line: (line.split(",")[1], 1))\
    .reduceByKey(lambda v1, v2: v1 + v2)
```

```

# Select the dates with at least 90% of the data centers associated with a high
power consumption
res1 = highPwrConsDCPerDay.filter(lambda t: t[1] >= threshold)\
                                .keys()

res1.saveAsTextFile(outputPath1)

#####
# PART 2
#####

# Consider the power consumptions and keep only the entries related to year 2021
# and obtain the following pairRDD
# key = codDC
# value = kWh
# and use a reduceByKey to sum the power consumption for the entire year for
# each data center

def mapCodDCpwrCons(line):
    fields = line.split(",")
    codDC = fields[0]
    pwrCons = float(fields[2])
    return (codDC, pwrCons)

yearlyPwrCons = pwrConsRDD.filter(lambda line:
line.split(",")[1].startswith("2021"))\
                        .map(mapCodDCpwrCons)\
                        .reduceByKey(lambda v1, v2: v1 + v2)

# for each data center, keep the continent information
# key = codDC
# value = continent
def mapCodDCContinent(line):
    fields = line.split(",")
    codDC = fields[0]
    continent = fields[4]
    return (codDC, continent)

dcAndContinent = dataCenterRDD.map(mapCodDCContinent)

# Join yearlyPwrCons with dcAndContinent and
# returns pairs
# key = continent

```

```

# value = (+1, kWhPerDataCenter2021)
continentOnePwr = yearlyPwrCons.join(dcAndContinent)\
    .map(lambda t: (t[1][1], (1, t[1][0])))

# Sum the value parts to compute for each continent
# the number of data centers and the total power consumption in the year 2021.
# key = continent
# value = (the number of data centers, avg power consumption in the year 2021)
#
# Finally, compute the avg power consumption for each continent
numDCAndAvgPwrCons = continentOnePwr\
    .reduceByKey(lambda t1, t2: (t1[0] + t2[0], t1[1] + t2[1]))\
    .mapValues(lambda t: (t[0], t[1]/t[0])).cache()

# compute the maximum number of data centers and the maximum avg consumption in
# the 2021 among the continents
maxDCAndConsumptionPerContinentThresholds = numDCAndAvgPwrCons\
    .values()\
    .reduce(lambda t1, t2: (max(t1[0],t2[0]), max(t1[1],t2[1])))

# select only those continents for which both constraints are satisfied
res2 = numDCAndAvgPwrCons\
    .filter(lambda t:
(t[1][0]==maxDCAndConsumptionPerContinentThresholds[0] and \
    t[1][1]==maxDCAndConsumptionPerContinentThresholds[
1]))\
    .keys()

res2.saveAsTextFile(outputPath2)

```