

15 FEB 2023 – HWC House Water Consumption

## MAPREDUCE & HADOOP

```
/**
 * Basic MapReduce Project - Mapper
 */
class Mapper1BigData extends Mapper<
    LongWritable,
    Text,
    Text,
    IntWritable> {

    protected void map(
        LongWritable key,
        Text value,
        Context context) throws IOException, InterruptedException {

        // HID,City,Country,YearBuilt
        String[] fields = value.toString().split(",");
        String city = fields[1];
        String country = fields[2];

        context.write(new Text(city + "," + country), new IntWritable(1));
    }
}

/**
 * Basic MapReduce Project - Reducer
 */
class Reducer1BigData extends Reducer<
    Text,
    IntWritable,
    Text,
    IntWritable> {

    @Override
    protected void reduce(
        Text key,
        Iterable<IntWritable> values,
        Context context) throws IOException, InterruptedException {
        int sum = 0;
        String[] fields = key.toString().split(",");
        String country = fields[1];
```

```

        for(IntWritable i : values)
            sum += i.get();
        context.write(new Text(country), new IntWritable(sum));
    }
}

/**
 * Basic MapReduce Project - Mapper
 */
class Mapper2BigData extends Mapper<
    Text,
    Text,
    Text,
    IntWritable> {

    protected void map(
        Text key,
        Text value,
        Context context) throws IOException, InterruptedException {

        int v = Integer.parseInt(value.toString());
        context.write(key, new IntWritable(v));
    }
}

/**
 * Basic MapReduce Project - Reducer
 */
class Reducer2BigData extends Reducer<
    Text,
    IntWritable,
    Text,
    DoubleWritable> {

    @Override
    protected void reduce(
        Text key,
        Iterable<IntWritable> values,
        Context context) throws IOException, InterruptedException {
        int sum = 0;
        int count = 0;

        for(IntWritable i : values) {

```

```
        sum += i.get();  
        count++;  
    }  
  
    double mean = ((double) sum) / count;  
    if(mean > 10000)  
        context.write(key, new DoubleWritable(mean));  
    }  
}
```

SPARK

```
from pyspark import SparkConf, SparkContext

conf = SparkConf().setAppName('Exam 15 feb 2023')
sc = SparkContext(conf = conf)

housePath = "data/Houses.txt"
consumptionPath = "data/MonthlyWaterConsumption.txt"

outputPath1 = "outPart1/"
outputPath2 = "outPart2/"

# Define the rdds associated with the input files

# HID, City, Country, YearBuilt
houseRDD = sc.textFile(housePath)

# HID, Month, M3
# Month in YYYY/MM format
consumptionRDD = sc.textFile(consumptionPath)

#####
# PART 1
#####

# filter water consumptions only for years 2022 and 2021
def filter22_21(l):
    fields = l.split(",")
    month = fields[1]

    return month.startswith("2021") or month.startswith("2022")

filteredCons = consumptionRDD.filter(filter22_21)

# compute <the trimesters for each record
# to do that, we use division between integers, with months starting from 0
(january = 0, february = 1 etc)
# month // 3 = trimesterID associated to the record
# example: january // 3 = 0, february // 3 = 0, march // 3 = 0, april // 3 = 1
etc.
# We compute then an RDD with
# key = HID, trimesterID, year
# value = water consumption
```

```
# and then sum all the values to compute the water consumed in that trimester
```

```
def HidTrimYear_Cons(l):  
    fields = l.split(",")  
    hid = fields[0]  
    yearMonth = fields[1]  
    consumption = float(fields[2])  
  
    f = yearMonth.split("/")  
    year = int(f[0])  
    month = int(f[1])  
    trimester = (month-1) // 3  
  
    return ((hid, trimester, year), consumption)
```

```
waterConsPerTrim = filteredCons\  
    .map(HidTrimYear_Cons)\  
    .reduceByKey(lambda v1, v2: v1 + v2)
```

```
# perform a map, obtaining  
# key = HID, trimester  
# value = (m3 in 2021, m3 in 2022)  
# where one of the two entries is 0  
# then, use a reduceByKey to sum the entries for each year (to obtain a single  
# entry per house and per trimester)
```

```
def hidTrim_cons2122(t):  
    hid = t[0][0]  
    trim = t[0][1]  
    year = t[0][2]  
    m3 = t[1]  
  
    if (year == 2021):  
        return ((hid, trim), (m3, 0.))  
    else:  
        return ((hid, trim), (0., m3))
```

```
houseConsPerTrim = waterConsPerTrim\  
    .map(hidTrim_cons2122)\  
    .reduceByKey(lambda v1, v2: (v1[0] + v2[0], v1[1] + v2[1]))
```

```
# select the pairs with an increasing consumption (m3 2022>m3 2021)
```

```

# and map to pairs with key = HID, value = +1
# and count for each house the number of increasing consumption associated to the
# same trimester
# and filter with count >= 3
housesWithIncreasingCons = houseConsPerTrim\
    .filter(lambda t: t[1][1] > t[1][0])\
    .map(lambda t: (t[0][0], 1))\
    .reduceByKey(lambda v1, v2: v1 + v2)\
    .filter(lambda t: t[1] >= 3)

# Prepare a pairRDD with
# key = hid
# value = city

def hid_city(l):
    fields = l.split(",")
    hid = fields[0]
    city = fields[1]

    return (hid, city)

houseCity = houseRDD.map(hid_city)

# join the two RDDs to compute result1
# map to HID, City
res1 = houseCity.join(housesWithIncreasingCons)\
    .map(lambda t: (t[0], t[1][0]))

res1.saveAsTextFile(outputPath1)

#####
# PART 2
#####

# compute for each year and houseID, the water consumption
# key = houseID, year
# value = water consumption

def HidYear_Cons(l):
    fields = l.split(",")
    hid = fields[0]
    yearMonth = fields[1]
    consumption = float(fields[2])

```

```

    f = yearMonth.split("/")
    year = int(f[0])

    return ((hid, year), consumption)

waterConsPerYearAndHouse = consumptionRDD\
    .map(HidYear_Cons)\
    .reduceByKey(lambda v1, v2: v1 + v2)

# Define windows of two consecutive year for each house
# flatMap each input to two pairs
#- key=(houseID, year ), value=( yearly consumption, +1)
#- key=(houseID, year+1), value=(-1*yearly consumption, +1)

def elements(p):
    pairs = []

    hid = p[0][0]
    year = p[0][1]

    AnnualCons= p[1]

    # - key=(houseID, year ), value=( yearly consumption, +1)
    pairs.append(((hid,year), (AnnualCons, 1)))

    # - key=(houseID, year+1), value=(-yearly consumption, +1)
    pairs.append(((hid,year+1), (-AnnualCons, 1)))

    return pairs

elementsWindows = waterConsPerYearAndHouse.flatMap(elements)

# Compute for each window the number of elements and the sum of the consumptions
# (each windows contains at most two values: one positive value for the current
year
# and a negative value for the previous year)
windowsElementsAndSumCons = elementsWindows\
    .reduceByKey(lambda p1, p2: (p1[0]+p2[0],
p1[1]+p2[1]))

# Select the windows with two elements and a negative sum of consumption
(decreasing consumption)

```

```

selectedWindows = windowsElementsAndSumCons\
    .filter(lambda p: p[1][1]==2 and p[1][0]<0)

# Map each window to (HID, +1) and remove duplicates.

def HID_1(p):
    hid = p[0][0]

    return (hid, 1)

houseAtLeastOneDecrease = selectedWindows\
    .map(HID_1)\
    .distinct()

# Join the previously computed RDD with houseCityRDD to get information related
# to the city
# and obtain the resulting RDD with
# key = city
# value = +1
# and count the number of houses for each city for which an annual consumption
# decrease was recorded
cityHouseConsDecr = houseAtLeastOneDecrease\
    .join(houseCity)\
    .map(lambda t: (t[1][1], 1))\
    .reduceByKey(lambda v1, v2: v1 + v2)

# filter and keep only the cities with count
# (number of houses with at least one annual decrease) > 2
citiesToDiscard = cityHouseConsDecr\
    .filter(lambda t:t[1] > 2)\
    .keys()

# Select the cities with at most two houses with at least one annual decrease
res2 = houseCity\
    .map(lambda p: p[1])\
    .distinct()\
    .subtract(citiesToDiscard)

res2.saveAsTextFile(outputPath2)

```